

Linear Discriminant Function - Passerini

Mattia Carolo - @Carolino96

Linear Discriminant Function

While for generative algorithms we infer knowledge from the modelled data distribution we have, with **discriminative learning** we focus directly on modelling the discriminant function. For example in classification with this method we directly model the decision boundary.

Formalization

Now we want to learn an $f(x)$ which maps x in y ($f : x \rightarrow y$). So we formalize the linear function like

$$f(x) = w^T x + w_0$$

The discriminant function is a linear combination of example features in which w_0 is the so called *bias* or *threshold* and it is the simplest possible discriminant function. The problem since it's linear we can't use it for complex problems but it's the most easy one to apply and once applied we can always expect the results so it's less prone to *overfitting*.

Linear Binary Classifier

Now a linear **binary** classifier it's a linear predictor which need to distinguish between two classes. So in this case we can get the normal linear discriminant function and get his sign so if the function will be positive the label will be 1 otherwise -1

$$f(x) = \text{sign}(w^T x + w_0)$$

if $f(x)$ is equal to 0 than we are on the **hyperplane** which will be the decision boundary. Moreover the weight vector w is orthogonal (\perp) to the decision

hyperplane. We can formalize it with a bit of algebra where we take two points on the hyperplane so their function is equal to zero as we do below where

$$\forall x, x' : f(x) = f(x') = 0$$

now if we replace them with their formalization we get that

$$w^T x + w_0 - w^T x' - w_0 = 0 \Rightarrow w^T (x - x') = 0$$

which means that w is orthogonal respect to $(x - x')$

Functional Margin

The functional margin can be seen as the confidence we have on a particular prediction. It's value is the value of $f(x)$ before applying the sign. The larger the value the more confident we are on that particular prediction.

Geometric margin

We call the geometric margin the distance between x and the hyperplane. We obtain it from the functional margin by doing

$$r^x = \frac{f(x)}{\|w\|}$$

where r^x will be the margin

we can see that the geometric margin is just the functional margin normalized

check demonstration on ot 27:17

Biological Motivation

shish

Perceptron

If we want to formalize a neuron we do it like a normal linear function made as

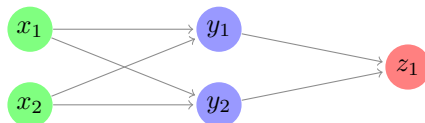
$$f(x) = \text{sign}(w^T x + w_0)$$

As we can see from the image below we have a set of inputs to the neuron from x_1 to x_m where every input will be multiplied for his weight(which can be positive

or negative). Later all this weighted inputs will be summed up together with w_0 and if the sum will be higher than the threshold we will predict 1 otherwise -1 (u need to surpass the threshold in order to activate it so that's why is called activation function)

add perceptron img

now if we take the function written before we can represent a series of primitive logical operations like AND, NAND or NOT since the results can be separated by a single line. In cases like the XOR in which we cannot learn the linear function itself since it doesn't exist, we can represent those cases by a network of two levels of perceptron.



For example in the case of the XOR we can decompose it in

$$x_1 \oplus x_2 = (x_1 \wedge \overline{x_2}) \vee (\overline{x_1} \wedge x_2)$$

where the node y_1 will take care of learning $(x_1 \wedge \overline{x_2})$, y_2 will learn $(\overline{x_1} \wedge x_2)$ and the node z_1 will learn the conjunction of the two.

Learning linear models

We already know the formula for the perceptron which is

$$f(x) = \text{sign}(w^T x)$$

but we need first a way to incorporate bias in our formula. To do this we use **augmented vectors** in which we add for both the feature vectors \hat{x} and the weight vector \hat{w} so they are made like

$$\hat{w} = \begin{bmatrix} w_0 \\ \mathbf{w} \end{bmatrix} \quad \hat{x} = \begin{bmatrix} 1 \\ \mathbf{x} \end{bmatrix}$$

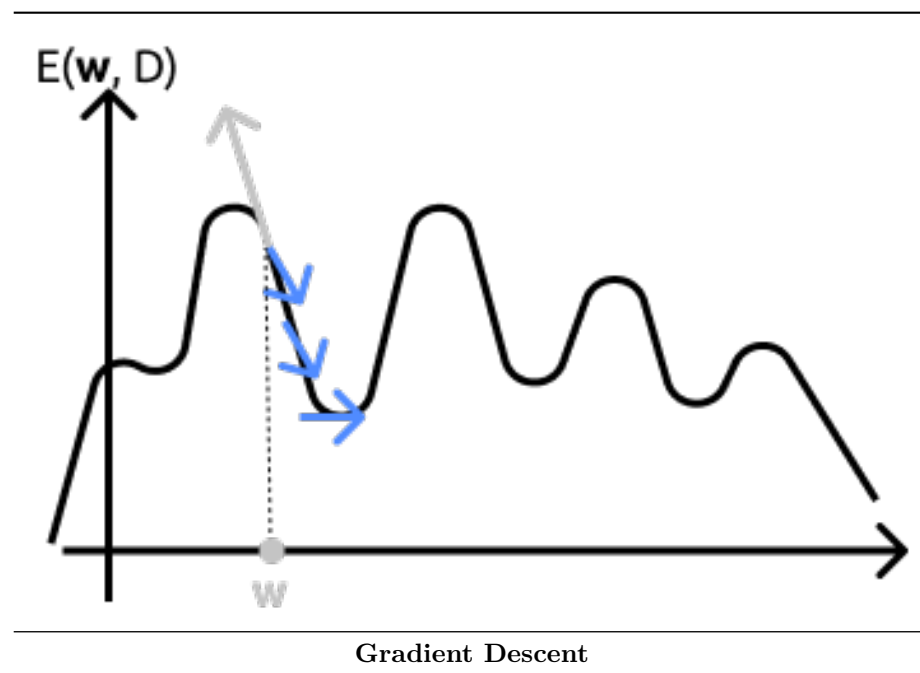
Now we need to find a function in order to maximize our parameters. One reasonable way to do it is measuring the estimate of the predictor given the ground truth and check whether it's correct or not like

$$E(\mathbf{w}, D) = \sum_{(\mathbf{x}, y) \in D} l(y, f(\mathbf{x}))$$

where: - $l(y, f(\mathbf{x}))$ measures the loss incurred on measuring $f(x)$ opposite to y
- D is our training set - l is the loss function - y the ground truth - $f(\mathbf{x})$ the prediction given

The problem in minimizing the error of the training set is that we risk that our model **overfits** the data which will result in a problem when we generalize it since other than the distribution the predictor learned even the noise from a particular set.

Gradient Descent



Let's take an error function $E(\mathbf{w}, D)$ with a single parameter w . Now we want to minimize this error. To do this we take the gradient of the error respect to W and we set it to zero which will result like

$$\nabla_{\mathbf{w}} E(\mathbf{w} : D) = 0$$

if the gradient reaches 0 it means we are either in a minimum or a maximum of the function. What we want is to reiterate this process until we reach a value approximable to zero. To do this first we set our \mathbf{w} . Now if we take the derivative (which is represented by the gray arrow) we can see that the result will be an upward arrow due to the gradient which follows the increase of a

function so since we want to minimize the error (**NOT** maximizing it) we'll take the opposite direction. Now to find the new \mathbf{w} we do

$$\mathbf{w} = \mathbf{w} - \eta \nabla_{\mathbf{w}} E$$

We can notice that we introduced the new term η which refers to the **learning rate** which tells us how much to go in that particular direction given by the gradient.

Values too large for η can result in a continuous oscillation around the minima without reaching it. Values too small from the other part will result in a slow learning or we could get stuck in a local minima. There are methods to set η on the fly

Parameter Learning

Now that we have a general understanding of the gradient descent we need to pick an error function to learn the parameter of our model. If we value accuracy we could say that a reasonable error function could be the **missclassification loss** $(1 - f(x))$ which is a very sensible measure of the classifier error but it's not good for gradient descent since in the function of the missclassification the only derivative can be zero or null.

Opposite to accuracy we can go for confidence where we use the **confidence loss**, the higher the confidence in the wrong prediction, the higher the result in the function: $E(\mathbf{w}, D) = \sum_{(\mathbf{x}, y) \in D_e} -yf(\mathbf{x})$. It's important to notice that we are

summing a subset of the training data. Infact D_e is just a subset of the current training set in which we took the examples where our classifier was taking wrong predictions. This measure is made like $yf(x) \leq 0$.

Now that we have an error function we can think about the gradient of the function. Applying the gradient we get that

$$\begin{aligned} \nabla E(\mathbf{w}, D) &= \nabla \sum_{(\mathbf{X}, y) \in D_E} -yf(x) \\ &= \nabla \sum_{(\mathbf{X}, y) \in D_E} -y(\mathbf{w}^T \mathbf{x}) \\ &= \nabla \sum_{(\mathbf{X}, y) \in D_E} -y\mathbf{x} \end{aligned}$$

so the update of \mathbf{w} at each iteration is $-\eta \nabla E(\mathbf{w}, D) = \eta \sum_{(\mathbf{x}, y) \in D_E} -y\mathbf{x}$ until we reach a zero gradient.

Still a problem rises. If we take a closer look we see that during gradient descent we iterate over all the examples which will result in slow computations in cases

where the dataset is really wide. This type of learning is called **batch learning** because everytime we compute the error on the batch of examples we are given

Stochastic training rule

An alternative is stochastic gradient descent. With this particular approach instead of computing the gradient of the error on all the batch, we make it stochastic by computing it only on a subset of the examples. First we initialize the weights randomly and then we iterate until all examples are correctly classified. When we meet a wrongly classified sample we immediately update the \mathbf{w} : $\mathbf{w} \leftarrow \mathbf{w} + \eta y \mathbf{x}$

Perceptron Regression

Let X be a matrix, in which we collect the training examples where $X \in \mathbb{R}^n \times \mathbb{R}^d$ with n rows, one for each sample ($n = |D|$), and d columns, one for each feature ($d = |\mathbf{x}|$). Let $\mathbf{y} \in \mathbb{R}^n$ be the output training vector where for every y_i is output for x_i .

At this point we can write linear regression as a set of linear equations where:

$$X\mathbf{w} = \mathbf{y}$$

in order to get the solution

$$\mathbf{w} = \mathbf{y}X^{-1}$$

The problem is that to reach to this particular solution we need X to be invertible which is not always our case. Usually the matrix X is a rectangular matrix with more rows than columns so we have more equations than unknowns so no exact solution typically exists (**overdetermined**).

Still this is not a big problem since we don't want to optimize a particular problem but to learn it.

Mean Squared Error

To address this problem we define an error function and we minimize it. The most common error function for regression is the **mean squared error** where the squared error is the difference between y and $f(\mathbf{x})$ where

$$E(\mathbf{w}, D) = \sum_{(\mathbf{x}, y) \in D} (y - f(\mathbf{x}))^2 = (\mathbf{y} - X\mathbf{w})^T (\mathbf{y} - X\mathbf{w})$$

For this formula it exists a closed form solution and can be used as a classification loss even if it's not the best. The solution goes like:

$$\begin{aligned}
\nabla E(\mathbf{w}; D) &= \nabla(\mathbf{y} - X\mathbf{w})^T(\mathbf{y} - X\mathbf{w}) \\
&= 2(\mathbf{y} - X\mathbf{w})^T(-X) = 0 \\
&= -2\mathbf{y}^T X + 2\mathbf{w}^T X^T X = 0 \\
\mathbf{w}^T X^T X &= \mathbf{y}^T X \\
X^T X \mathbf{w} &= X^T \mathbf{y} \\
\mathbf{w} &= (X^T X)^{-1} X^T \mathbf{y}
\end{aligned}$$

where in particular: - At the second step if we zero we get the third term - Now at the 4th step we get that w is transposed but we need the non transposed value so we transpose all the term the holds w where $(\mathbf{w}^T X^T X)$ becomes $(\mathbf{w}^T)^T (X^T X)^T$ where \mathbf{w} will return to the normal value while the second term will become $(X^T X)$

Check cheat sheet for further explanation on the solution